

Install Kubernetes Cluster on Ubuntu 20.04 with kubeadm

By [Josphat Mutai](#) - November 30, 2021

Kubernetes is a tool for orchestrating and managing containerized applications at scale on on-premise server or across hybrid cloud environments. Kubeadm is a tool provided with Kubernetes to help users install a production ready Kubernetes cluster with best practices enforcement. This tutorial will demonstrate how one can install a Kubernetes Cluster on Ubuntu 20.04 with kubeadm.

For Debian installation: [Deploy Kubernetes Cluster on Debian 10 with Kubespray](#)

For Rocky Linux 8: [Install Kubernetes Cluster on Rocky Linux 8 with Kubeadm & CRI-O](#)

There are two server types used in deployment of Kubernetes clusters:

- **Master:** A Kubernetes Master is where control API calls for the pods, replications controllers, services, nodes and other components of a Kubernetes cluster are executed



The minimum requirements for the viable setup are:

- Memory: **2 GiB** or more of RAM per machine
- CPUs: At least **2 CPUs** on the control plane machine.
- Internet connectivity for pulling containers required (Private registry can also be used)
- Full network connectivity between machines in the cluster – This is private or public

Install Kubernetes Cluster on Ubuntu 20.04

My Lab setup contain three servers. One control plane machine and two nodes to be used for running containerized workloads. You can add more nodes to suit your desired use case and load, for example using **three** control plane nodes for HA.

Server Type	Server Hostname	Specs
Master	k8s-master01.computingforgeeks.com	4GB Ram, 2vcpus
Worker	k8s-worker01.computingforgeeks.com	4GB Ram, 2vcpus
Worker	k8s-worker02.computingforgeeks.com	4GB Ram, 2vcpus

Step 1: Install Kubernetes Servers

Provision the servers to be used in the deployment of Kubernetes on Ubuntu 20.04. The setup process will vary depending on the virtualization or cloud environment you're using.

Once the servers are ready, update them.

```
sudo apt update
sudo apt -y upgrade && sudo systemctl reboot
```

Step 2: Install kubelet, kubeadm and kubectl

Once the servers are rebooted, add Kubernetes repository for Ubuntu 20.04 to all the servers.

```
sudo apt update
```



```
echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" | sudo  
tee /etc/apt/sources.list.d/kubernetes.list
```

Then install required packages.

```
sudo apt update  
sudo apt -y install vim git curl wget kubelet kubeadm kubectl  
sudo apt-mark hold kubelet kubeadm kubectl
```

Confirm installation by checking the version of kubectl.

```
$ kubectl version --client && kubeadm version  
Client Version: version.Info{Major:"1", Minor:"22",  
GitVersion:"v1.22.2",  
GitCommit:"8b5a19147530eaac9476b0ab82980b4088bbc1b2",  
GitTreeState:"clean", BuildDate:"2021-09-15T21:38:50Z",  
GoVersion:"go1.16.8", Compiler:"gc", Platform:"linux/amd64"}  
kubeadm version: &version.Info{Major:"1", Minor:"22",  
GitVersion:"v1.22.2",  
GitCommit:"8b5a19147530eaac9476b0ab82980b4088bbc1b2",  
GitTreeState:"clean", BuildDate:"2021-09-15T21:37:34Z",  
GoVersion:"go1.16.8", Compiler:"gc", Platform:"linux/amd64"}
```

Step 3: Disable Swap

Turn off swap.

```
sudo sed -i '/ swap / s/^(.*)$/#\1/g' /etc/fstab  
sudo swapoff -a
```

Enable kernel modules and configure sysctl.

```
sudo modprobe overlay
sudo modprobe br_netfilter
```

Add some settings to sysctl

```
sudo tee /etc/sysctl.d/kubernetes.conf<<EOF
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
EOF
```

Reload sysctl

```
sudo sysctl --system
```

Step 4: Install Container runtime

To run containers in Pods, Kubernetes uses a container runtime. Supported container runtimes are:

- Docker
- CRI-O
- Containerd

NOTE: You have to choose one runtime at a time.

Installing Docker runtime:

Add repo and Install packages

```
sudo apt update
sudo apt install -y curl gnupg2 software-properties-common apt-
transport-https ca-certificates
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-
key add -
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com
/linux/ubuntu $(lsb_release -cs) stable"
sudo apt update
sudo apt install -y containerd.io docker-ce docker-ce-cli
```

Create required directories

```
sudo mkdir -p /etc/systemd/system/docker.service.d
```

Create daemon json config file

```
sudo tee /etc/docker/daemon.json <<EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
```

```
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```

Start and enable Services

```
sudo systemctl daemon-reload
sudo systemctl restart docker
sudo systemctl enable docker
```

Installing CRI-O:

Ensure you load modules

```
sudo modprobe overlay
sudo modprobe br_netfilter
```

Set up required sysctl params

```
sudo tee /etc/sysctl.d/kubernetes.conf<<EOF
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
EOF
```

Reload sysctl

```
sudo sysctl --system
```

Add Cri-o repo

```
sudo su -
OS="xUbuntu_20.04"
VERSION=1.22
echo "deb https://download.opensuse.org/repositories/devel:/kubic:
/libcontainers:/stable/$OS/ /" > /etc/apt/sources.list.d
/devel:kubic:libcontainers:stable.list
echo "deb http://download.opensuse.org/repositories/devel:/kubic:
/libcontainers:/stable:/cri-o:/$VERSION/$OS/ /" > /etc/apt
/sources.list.d/devel:kubic:libcontainers:stable:cri-o:$VERSION.list
curl -L https://download.opensuse.org/repositories
/devel:kubic:libcontainers:stable:cri-o:$VERSION/$OS/Release.key |
apt-key add -
curl -L https://download.opensuse.org/repositories/devel:/kubic:
/libcontainers:/stable/$OS/Release.key | apt-key add -
```

Install CRI-O

```
sudo apt update
sudo apt install cri-o cri-o-runc
```

Start and enable Service

```
sudo systemctl daemon-reload
sudo systemctl restart crio
sudo systemctl enable crio
sudo systemctl status crio
```

Installing Containerd:*# Configure persistent loading of modules*

```
sudo tee /etc/modules-load.d/containerd.conf <<EOF
overlay
br_netfilter
EOF
```

Load at runtime

```
sudo modprobe overlay
sudo modprobe br_netfilter
```

Ensure sysctl params are set

```
sudo tee /etc/sysctl.d/kubernetes.conf<<EOF
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
EOF
```

Reload configs

```
sudo sysctl --system
```

Install required packages

```
sudo apt install -y curl gnupg2 software-properties-common apt-
transport-https ca-certificates
```

Add Docker repo

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-
key add -
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com
/linux/ubuntu $(lsb_release -cs) stable"
```

```
# Configure containerd and start service
sudo su -
mkdir -p /etc/containerd
containerd config default > /etc/containerd/config.toml
```

```
# restart containerd
sudo systemctl restart containerd
sudo systemctl enable containerd
systemctl status containerd
```

To use the systemd cgroup driver, set **plugins.cri.systemd_cgroup = true** in `/etc/containerd/config.toml` . When using kubeadm, manually configure the [cgroup driver for kubelet](#)

Step 5: Initialize master node

Login to the server to be used as master and make sure that the `br_netfilter` module is loaded:

```
$ lsmod | grep br_netfilter
br_netfilter      22256  0
bridge           151336  2 br_netfilter,ebtable_broute
```

Enable kubelet service.

```
sudo systemctl enable kubelet
```

We now want to initialize the machine that will run the control plane components which includes `etcd` (the cluster database) and the API Server.

```
[config/images] Pulled k8s.gcr.io/kube-apiserver:v1.22.2
[config/images] Pulled k8s.gcr.io/kube-controller-manager:v1.22.2
[config/images] Pulled k8s.gcr.io/kube-scheduler:v1.22.2
[config/images] Pulled k8s.gcr.io/kube-proxy:v1.22.2
[config/images] Pulled k8s.gcr.io/pause:3.5
[config/images] Pulled k8s.gcr.io/etcd:3.5.0-0
[config/images] Pulled k8s.gcr.io/coredns/coredns:v1.8.4
```

If you have multiple CRI sockets, please use `--cri-socket` to select one:

CRI-O

```
sudo kubeadm config images pull --cri-socket /var/run/crio/crio.sock
```

Containerd

```
sudo kubeadm config images pull --cri-socket /run/containerd
/containerd.sock
```

Docker

```
sudo kubeadm config images pull --cri-socket /var/run/docker.sock
```

These are the basic `kubeadm init` options that are used to bootstrap cluster.

- control-plane-endpoint** : set the shared endpoint for all control-plane nodes. Can be DNS/IP
- pod-network-cidr** : Used to set a Pod network add-on CIDR
- cri-socket** : Use if have more than one container runtime to set runtime socket path
- apiserver-advertise-address** : Set advertise address for this particular control-plane node's API server

Bootstrap without shared endpoint

To bootstrap a cluster without using DNS endpoint, run:

```
sudo kubeadm init \
  --pod-network-cidr=192.168.0.0/16
```

Bootstrap with shared endpoint (DNS name for control plane API)

Set cluster endpoint DNS name or add record to `/etc/hosts` file.

```
$ sudo vim /etc/hosts
172.29.20.5 k8s-cluster.computingforgeeks.com
```

Create cluster:


```
--upload-certs \  
--control-plane-endpoint=k8s-cluster.computingforgeeks.com
```

Note: If *192.168.0.0/16* is already in use within your network you must select a different pod network CIDR, replacing *192.168.0.0/16* in the above command.

Container runtime sockets:

Runtime	Path to Unix domain socket
Docker	/var/run/docker.sock
containerd	/run/containerd/containerd.sock
CRI-O	/var/run/crio/crio.sock

You can optionally pass Socket file for runtime and advertise address depending on your setup.

CRI-O

```
sudo kubeadm init \  
--pod-network-cidr=192.168.0.0/16 \  
--cri-socket /var/run/crio/crio.sock \  
--upload-certs \  
--control-plane-endpoint=k8s-cluster.computingforgeeks.com
```

Containerd

```
sudo kubeadm init \  
--pod-network-cidr=192.168.0.0/16 \  
--cri-socket /run/containerd/containerd.sock \  
--upload-certs \  
--control-plane-endpoint=k8s-cluster.computingforgeeks.com
```

```
--pod-network-cidr=192.168.0.0/16 \  
--cri-socket /var/run/docker.sock \  
--upload-certs \  
--control-plane-endpoint=k8s-cluster.computingforgeeks.com
```

Here is the output of my initialization command.

```
....  
[init] Using Kubernetes version: v1.22.2  
[preflight] Running pre-flight checks  
      [WARNING Firewall]: firewalld is active, please ensure ports  
[6443 10250] are open or your cluster may not function correctly  
[preflight] Pulling images required for setting up a Kubernetes  
cluster  
[preflight] This might take a minute or two, depending on the speed  
of your internet connection  
[preflight] You can also perform this action in beforehand using  
'kubeadm config images pull'  
[kubelet-start] Writing kubelet environment file with flags to file  
"/var/lib/kubelet/kubeadm-flags.env"  
[kubelet-start] Writing kubelet configuration to file "/var/lib  
/kubelet/config.yaml"  
[kubelet-start] Starting the kubelet  
[certs] Using certificateDir folder "/etc/kubernetes/pki"  
[certs] Using existing ca certificate authority  
[certs] Using existing apiserver certificate and key on disk  
[certs] Using existing apiserver-kubelet-client certificate and key  
on disk  
[certs] Using existing front-proxy-ca certificate authority  
[certs] Using existing front-proxy-client certificate and key on disk
```

```
[certs] Using existing etcd/peer certificate and key on disk
[certs] Using existing etcd/healthcheck-client certificate and key on
disk
[certs] Using existing apiserver-etcd-client certificate and key on
disk
[certs] Using the existing "sa" key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Using existing kubeconfig file: "/etc/kubernetes
/admin.conf"
[kubeconfig] Using existing kubeconfig file: "/etc/kubernetes
/kubelet.conf"
[kubeconfig] Using existing kubeconfig file: "/etc/kubernetes
/controller-manager.conf"
[kubeconfig] Using existing kubeconfig file: "/etc/kubernetes
/scheduler.conf"
[control-plane] Using manifest folder "/etc/kubernetes/manifests"
[control-plane] Creating static Pod manifest for "kube-apiserver"
[control-plane] Creating static Pod manifest for "kube-controller-
manager"
W0611 22:34:23.276374    4726 manifests.go:225] the default kube-
apiserver authorization-mode is "Node,RBAC"; using "Node,RBAC"
[control-plane] Creating static Pod manifest for "kube-scheduler"
W0611 22:34:23.278380    4726 manifests.go:225] the default kube-
apiserver authorization-mode is "Node,RBAC"; using "Node,RBAC"
[etcd] Creating static Pod manifest for local etcd in
"/etc/kubernetes/manifests"
[wait-control-plane] Waiting for the kubelet to boot up the control
plane as static Pods from directory "/etc/kubernetes/manifests". This
can take up to 4m0s
[apiclient] All control plane components are healthy after 8.008181
seconds
[upload-config] Storing the configuration used in ConfigMap "kubeadm-
config" in the "kube-system" Namespace
[kubelet] Creating a ConfigMap "kubelet-config-1.21" in namespace
kube-system with the configuration for the kubelets in the cluster
[upload-certs] Skipping phase. Please see --upload-certs
[mark-control-plane] Marking the node k8s-
master01.computingforgeeks.com as control-plane by adding the label
"node-role.kubernetes.io/master=''"
[mark-control-plane] Marking the node k8s-
master01.computingforgeeks.com as control-plane by adding the taints
[node-role.kubernetes.io/master:NoSchedule]
[bootstrap-token] Using token: zoy8cq.6v349sx9ass8dzyj
```

```
tokens to get nodes
[bootstrap-token] configured RBAC rules to allow Node Bootstrap
tokens to post CSRs in order for nodes to get long term certificate
credentials
[bootstrap-token] configured RBAC rules to allow the csrapprover
controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] configured RBAC rules to allow certificate rotation
for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-
public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point
to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

You can now join any number of control-plane nodes by copying certificate authorities and service account keys on each node and then running the following as root:

```
kubeadm join k8s-cluster.computingforgeeks.com:6443 --token
sr4l2l.2kvot0pfalh5o4ik \
    --discovery-token-ca-cert-hash
sha256:c692fb047e15883b575bd6710779dc2c5af8073f7cab460abd181fd3ddb29a
18 \
    --control-plane
```

Then you can join any number of worker nodes by running the following on each as root:

```
--discovery-token-ca-cert-hash  
sha256:c692fb047e15883b575bd6710779dc2c5af8073f7cab460abd181fd3ddb29a  
18
```

Configure kubectl using commands in the output:

```
mkdir -p $HOME/.kube  
sudo cp -f /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Check cluster status:

```
$ kubectl cluster-info  
Kubernetes master is running at https://k8s-  
cluster.computingforgeeks.com:6443  
KubeDNS is running at https://k8s-cluster.computingforgeeks.com:6443  
/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
```

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

Additional Master nodes can be added using the command in installation output:

```
kubeadm join k8s-cluster.computingforgeeks.com:6443 --token  
sr4l2l.2kvot0pfalh5o4ik \  
--discovery-token-ca-cert-hash  
sha256:c692fb047e15883b575bd6710779dc2c5af8073f7cab460abd181fd3ddb29a  
18 \  
--control-plane
```

```
kubectl create -f https://docs.projectcalico.org/manifests/tigera-  
operator.yaml  
kubectl create -f https://docs.projectcalico.org/manifests/custom-  
resources.yaml
```

You should see the following output.

```
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.p  
rojectcalico.org created  
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcal  
ico.org created  
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.pro  
jectcalico.org created  
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd  
.projectcalico.org created  
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd  
.projectcalico.org created  
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.c  
rd.projectcalico.org created  
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.p  
rojectcalico.org created  
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.proje  
ctcalico.org created  
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectc  
alico.org created  
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.project  
calico.org created  
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.project  
calico.org created  
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcali  
co.org created  
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigur  
ations.crd.projectcalico.org created  
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.pro  
jectcalico.org created
```

```

customresourcedefinition.apiextensions.k8s.io/apiservers.operator.tig
era.io created
customresourcedefinition.apiextensions.k8s.io/imagesets.operator.tige
ra.io created
customresourcedefinition.apiextensions.k8s.io/installations.operator.
tigera.io created
customresourcedefinition.apiextensions.k8s.io/tigerastatuses.operator
.tigera.io created
namespace/tigera-operator created
Warning: policy/v1beta1 PodSecurityPolicy is deprecated in v1.21+,
unavailable in v1.25+
podsecuritypolicy.policy/tigera-operator created
serviceaccount/tigera-operator created
clusterrole.rbac.authorization.k8s.io/tigera-operator created
clusterrolebinding.rbac.authorization.k8s.io/tigera-operator created
deployment.apps/tigera-operator created
.....
installation.operator.tigera.io/default created
apiserver.operator.tigera.io/default created

```

Confirm that all of the pods are running:

```

$ watch kubectl get pods --all-namespaces

```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	calico-kube-controllers-76d4774d89-nfqrr	1/1	Running	0	2m52s
kube-system	calico-node-kpprr	1/1	Running	0	2m52s
kube-system	coredns-66bff467f8-9bxgm	1/1	Running	0	7m43s
kube-system	coredns-66bff467f8-jgwln	1/1	Running	0	7m43s
kube-system	etcd-k8s-master01.computingforgeeks.com	1/1	Running	0	7m58s



```

kube-system    kube-controller-manager-k8s-
master01.computingforgeeks.com    1/1    Running    0            7m58s
kube-system    kube-proxy-bt7ff
1/1    Running    0            7m43s
kube-system    kube-scheduler-k8s-master01.computingforgeeks.com
1/1    Running    0            7m58s

```

Confirm master node is ready:

CRI-O

\$ *kubectl get nodes -o wide*

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP
EXTERNAL-IP	OS-IMAGE		KERNEL-VERSION		CONTAINER-
					RUNTIME
ubuntu	Ready	control-plane,master	38s	v1.22.2	
143.198.114.46	<none>		Ubuntu 20.04.3 LTS	5.4.0-88-generic	
					cri-o://1.22.0

Containerd

\$ *kubectl get nodes -o wide*

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP
EXTERNAL-IP	OS-IMAGE		KERNEL-VERSION		CONTAINER-
					RUNTIME
ubuntu	Ready	control-plane,master	15m	v1.22.2	
143.198.114.46	<none>		Ubuntu 20.04.3 LTS	5.4.0-88-generic	
					containerd://1.4.11

Docker

\$ *kubectl get nodes -o wide*

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP
EXTERNAL-IP	OS-IMAGE		KERNEL-VERSION		CONTAINER-RUNTIME
k8s-master01	Ready	master	64m	v1.22.2	135.181.28.113
<none>	Ubuntu 20.04	LTS	5.4.0-37-generic		docker://20.10.8

Step 7: Add worker nodes

With the control plane ready you can add worker nodes to the cluster for running scheduled workloads.

If endpoint address is not in DNS, add record to `/etc/hosts`.

```
$ sudo vim /etc/hosts
172.29.20.5 k8s-cluster.computingforgeeks.com
```

The join command that was given is used to add a worker node to the cluster.

```
kubeadm join k8s-cluster.computingforgeeks.com:6443 \
  --token sr4l2l.2kvot0pfalh5o4ik \
  --discovery-token-ca-cert-hash
sha256:c692fb047e15883b575bd6710779dc2c5af8073f7cab460abd181fd3ddb29a
18
```

Output:

```
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n
kube-system get cm kubeadm-config -oyaml'
[kubelet-start] Downloading configuration for the kubelet from the
"kubelet-config-1.21" ConfigMap in the kube-system namespace
[kubelet-start] Writing kubelet configuration to file "/var/lib
/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file
"/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS
Bootstrap...
```

This node has joined the cluster:

```
* Certificate signing request was sent to apiserver and a response
```

Run below command on the control-plane to see if the node joined the cluster.

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8s-master01.computingforgeeks.com	Ready	master	10m	v1.22.2
k8s-worker01.computingforgeeks.com	Ready	<none>	50s	v1.22.2
k8s-worker02.computingforgeeks.com	Ready	<none>	12s	v1.22.2

```
$ kubectl get nodes -o wide
```

If the join token is expired, refer to our guide on how to join worker nodes.

- [Join new Kubernetes Worker Node to an existing Cluster](#)

Step 8: Deploy application on cluster

If you only have a single node cluster, check our guide on how to run container pods on master nodes:

- [Scheduling Pods on Kubernetes Control plane \(Master\) Nodes](#)

We need to validate that our cluster is working by deploying an application.

```
kubectl apply -f https://k8s.io/examples/pods/commands.yaml
```

Check to see if pod started



```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
command-demo	0/1	Completed	0	16s

Step 9: Install Kubernetes Dashboard (Optional)

Kubernetes dashboard can be used to deploy containerized applications to a Kubernetes cluster, troubleshoot your containerized application, and manage the cluster resources.

Refer to our guide for installation: [How To Install Kubernetes Dashboard with NodePort](#)

Step 10: Install Metrics Server (For checking Pods and Nodes resource usage)

Metrics Server is a cluster-wide aggregator of resource usage data. It collects metrics from the *Summary API*, exposed by **Kubelet** on each node. Use our guide below to deploy it:

- [How To Deploy Metrics Server to Kubernetes Cluster](#)

Step 11: Deploy Prometheus / Grafana Monitoring

Prometheus is a full fledged solution that enables you to access advanced metrics capabilities in a Kubernetes cluster. Grafana is used for analytics and interactive visualization of metrics that's collected and stored in Prometheus database. We have a complete guide on how to setup complete monitoring stack on Kubernetes Cluster:

- [Setup Prometheus and Grafana on Kubernetes using prometheus-operator](#)

Step 12: Persistent Storage Configuration ideas (Optional)

If you're also looking for a Persistent storage solution for your Kubernetes, checkout:

- [How To Deploy Rook Ceph Storage on Kubernetes Cluster](#)
- [Ceph Persistent Storage for Kubernetes with Cephfs](#)
- [Persistent Storage for Kubernetes with Ceph RBD](#)
- [How To Configure Kubernetes Dynamic Volume Provisioning With Heketi & GlusterFS](#)

If Nginx is your preferred Ingress controller for Kubernetes workloads, you can use our guide for the installation process:

- [Deploy Nginx Ingress Controller on Kubernetes using Helm Chart](#)

More guides:

- [Using Horizontal Pod Autoscaler on Kubernetes Cluster](#)
- [Install Kubernetes Metrics Server](#)

Similar Kubernetes deployment guides:

- [Install Kubernetes Cluster on CentOS 7 with kubeadm](#)
- [Install Production Kubernetes Cluster with Rancher RKE](#)
- [How To Deploy Lightweight Kubernetes Cluster in 5 minutes with K3s](#)
- [Deploy Production Ready Kubernetes Cluster with Ansible & Kubespray](#)

**Your support is our everlasting motivation,
that cup of coffee is what keeps us going!**

As we continue to grow, we would wish to reach and impact more people who visit and take advantage of the guides we have on our blog. This is a big task for us and we are so far extremely grateful for the kind people who have shown amazing support for our work over the time we have been online.

Thank You for your support as we work to give you the best of guides and articles. Click below to buy us a coffee.



Josphat Mutai

<https://computingforgeeks.com/>

Founder of Computingforgeeks. Expertise in Virtualization, Cloud, Linux/UNIX Administration, Automation, Storage Systems, Containers, Server Clustering e.t.c.



report this ad

